

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: data = pd.read_csv('bank-full (1).csv', sep=';')
data.head()
```

```
Out[2]:
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	poutcom
0	58	management	married	tertiary	no	2143	yes	no	unknown	5	may	261	1	-1	0	unknow
1	44	technician	single	secondary	no	29	yes	no	unknown	5	may	151	1	-1	0	unknow
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	may	76	1	-1	0	unknow
3	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	may	92	1	-1	0	unknow
4	33	unknown	single	unknown	no	1	no	no	unknown	5	may	198	1	-1	0	unknow

```
In [3]: data.shape
```

```
Out[3]: (45211, 17)
```

```
In [4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         45211 non-null  int64
1   job         45211 non-null  object
2   marital     45211 non-null  object
3   education   45211 non-null  object
4   default     45211 non-null  object
5   balance     45211 non-null  int64
6   housing     45211 non-null  object
7   loan        45211 non-null  object
8   contact     45211 non-null  object
9   day         45211 non-null  int64
10  month       45211 non-null  object
11  duration    45211 non-null  int64
12  campaign    45211 non-null  int64
13  pdays       45211 non-null  int64
14  previous    45211 non-null  int64
15  poutcome    45211 non-null  object
16  y           45211 non-null  object
dtypes: int64(7), object(10)
memory usage: 5.9+ MB
```


```
In [5]: data.isnull().sum()
```

```
Out[5]: age         0
job         0
marital     0
education   0
default     0
balance     0
housing     0
loan        0
contact     0
day         0
month       0
duration    0
campaign    0
pdays      0
previous    0
poutcome    0
y           0
dtype: int64
```

```
In [6]: import seaborn as sns
```

```
In [7]: sns.heatmap(data.isnull(),cmap='Blues')
# there are no null values
```

Out[7]: <AxesSubplot:>



```
In [8]: data.dtypes
```

```
Out[8]: age          int64
job          object
marital      object
education    object
default      object
balance      int64
housing      object
loan         object
contact      object
day          int64
month        object
duration     int64
campaign     int64
pdays       int64
previous     int64
poutcome     object
y            object
dtype: object
```

```
In [9]: # there are no duplicates
data.duplicated()
```

```
Out[9]: 0      False
1      False
2      False
3      False
4      False
...
45206   False
45207   False
45208   False
45209   False
45210   False
Length: 45211, dtype: bool
```

```
In [10]: data.head(2)
```

```
Out[10]:
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	poutcom
0	58	management	married	tertiary	no	2143	yes	no	unknown	5	may	261	1	-1	0	unknow
1	44	technician	single	secondary	no	29	yes	no	unknown	5	may	151	1	-1	0	unknow

```
In [11]: # checking the categorical values
categorical = [var for var in data.columns if data[var].dtype=='O']

print('There are {} categorical variables\n'.format(len(categorical)))
```

```
print('The categorical variables are :', categorical)
```

There are 10 categorical variables

The categorical variables are : ['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'poutcome', 'y']

```
In [12]: data[categorical].isnull().sum()
```

```
Out[12]: job          0
marital        0
education      0
default        0
housing        0
loan           0
contact        0
month          0
poutcome       0
y             0
dtype: int64
```

```
In [13]: # covert the categorical into dummies
data = pd.get_dummies(data,columns=['job','marital','education','contact','month','poutcome'])
```

```
In [14]: data.head()
```

```
Out[14]:
```

	age	default	balance	housing	loan	day	duration	campaign	pdays	previous	...	month_jun	month_mar	month_may	month_nov	month_oct
0	58	no	2143	yes	no	5	261	1	-1	0	...	0	0	1	0	0
1	44	no	29	yes	no	5	151	1	-1	0	...	0	0	1	0	0
2	33	no	2	yes	yes	5	76	1	-1	0	...	0	0	1	0	0
3	47	no	1506	yes	no	5	92	1	-1	0	...	0	0	1	0	0
4	33	no	1	no	no	5	198	1	-1	0	...	0	0	1	0	0

5 rows × 49 columns

```
In [15]: # convert the binary to string
data['default'] = np.where(data['default'].str.contains('yes'),1,0)
data['housing'] = np.where(data['housing'].str.contains('yes'),1,0)
data['loan'] = np.where(data['loan'].str.contains('yes'),1,0)
data['y'] = np.where(data['y'].str.contains('yes'),1,0)
data
```

```
Out[15]:
```

	age	default	balance	housing	loan	day	duration	campaign	pdays	previous	...	month_jun	month_mar	month_may	month_nov	month_oct
0	58	0	2143	1	0	5	261	1	-1	0	...	0	0	1	0	0
1	44	0	29	1	0	5	151	1	-1	0	...	0	0	1	0	0
2	33	0	2	1	1	5	76	1	-1	0	...	0	0	1	0	0
3	47	0	1506	1	0	5	92	1	-1	0	...	0	0	1	0	0
4	33	0	1	0	0	5	198	1	-1	0	...	0	0	1	0	0
...
45206	51	0	825	0	0	17	977	3	-1	0	...	0	0	0	1	0
45207	71	0	1729	0	0	17	456	2	-1	0	...	0	0	0	1	0
45208	72	0	5715	0	0	17	1127	5	184	3	...	0	0	0	1	0
45209	57	0	668	0	0	17	508	4	-1	0	...	0	0	0	1	0
45210	37	0	2971	0	0	17	361	2	188	11	...	0	0	0	1	0

45211 rows × 49 columns

```
In [16]: data.iloc[:,15:30]
```

```
Out[16]:
```

	job_management	job_retired	job_self-employed	job_services	job_student	job_technician	job_unemployed	job_unknown	marital_divorced	marital_married
--	----------------	-------------	-------------------	--------------	-------------	----------------	----------------	-------------	------------------	-----------------

0	1	0	0	0	0	0	0	0	0
1	0	0	0	0	0	1	0	0	0
2	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	1	0
...
45206	0	0	0	0	0	1	0	0	0
45207	0	1	0	0	0	0	0	0	1
45208	0	1	0	0	0	0	0	0	0
45209	0	0	0	0	0	0	0	0	0
45210	0	0	0	0	0	0	0	0	0

```
In [17]: X = data.drop(columns='y')
X
```

Out[17]:	age	default	balance	housing	loan	day	duration	campaign	pdays	previous	...	month_jun	month_mar	month_may	month_nov	n
0	58	0	2143	1	0	5	261	1	-1	0	...	0	0	1	0	
1	44	0	29	1	0	5	151	1	-1	0	...	0	0	1	0	
2	33	0	2	1	1	5	76	1	-1	0	...	0	0	1	0	
3	47	0	1506	1	0	5	92	1	-1	0	...	0	0	1	0	
4	33	0	1	0	0	5	198	1	-1	0	...	0	0	1	0	
...	
45206	51	0	825	0	0	17	977	3	-1	0	...	0	0	0	1	
45207	71	0	1729	0	0	17	456	2	-1	0	...	0	0	0	1	
45208	72	0	5715	0	0	17	1127	5	184	3	...	0	0	0	1	
45209	57	0	668	0	0	17	508	4	-1	0	...	0	0	0	1	
45210	37	0	2971	0	0	17	361	2	188	11	...	0	0	0	1	

```
In [18]: y_data = data.loc[:, 'y']  
y_data
```

```
Out[18]:
```

0	0
1	0
2	0
3	0
4	0
	..
45206	1
45207	1
45208	1
45209	0
45210	0

Name: y, Length: 45211, dtype: int32

```
In [19]: y_data.unique()
```

```
Out[19]: array([0, 1])
```

```
In [20]: from sklearn.model_selection import train_test_split
```

```
In [21]: X_train, X_test, y_train, y_test = train_test_split(X,y_data,test_size=0.3)
X_train, X_test, y_train, y_test
```

```
Out[21]:
```

	age	default	balance	housing	loan	day	duration	campaign	pdays
44745	62	0	2801	1	1	9	261	1	183
44864	46	0	7485	0	0	23	145	1	779

34749	34	0	649	1	0	6	224	1	363
18099	42	0	2613	1	0	30	174	6	-1
1622	32	0	536	1	0	9	208	2	-1
...
8120	60	0	-79	1	0	2	49	1	-1
39381	54	0	1188	1	0	22	89	1	-1
18512	35	0	-195	0	0	31	309	4	-1
29574	33	0	1883	1	0	3	121	5	256
655	33	0	-349	1	0	6	191	1	-1

	previous	...	month_jun	month_mar	month_may	month_nov	month_oct	\
44745	1	...	0	0	0	0	0	0
44864	2	...	0	0	0	0	0	0
34749	1	...	0	0	1	0	0	0
18099	0	...	0	0	0	0	0	0
1622	0	...	0	0	1	0	0	0
...
8120	0	...	1	0	0	0	0	0
39381	0	...	0	0	1	0	0	0
18512	0	...	0	0	0	0	0	0
29574	1	...	0	0	0	0	0	0
655	0	...	0	0	1	0	0	0

	month_sep	poutcome_failure	poutcome_other	poutcome_success	\
44745	1	0	0	1	
44864	1	1	0	0	
34749	0	0	1	0	
18099	0	0	0	0	
1622	0	0	0	0	
...	
8120	0	0	0	0	
39381	0	0	0	0	
18512	0	0	0	0	
29574	0	1	0	0	
655	0	0	0	0	

	poutcome_unknown
44745	0
44864	0
34749	0
18099	1
1622	1
...	...
8120	1
39381	1
18512	1
29574	0
655	1

[31647 rows x 48 columns],

	age	default	balance	housing	loan	day	duration	campaign	pdays	\
10547	42	0	167	0	0	16	119	1	-1	
40059	33	0	506	0	0	4	176	1	91	
26737	32	0	6982	1	0	20	224	2	183	
21071	57	0	209	0	0	14	56	4	-1	
39637	29	0	2907	1	0	26	150	1	-1	
...	
38989	31	0	1374	1	0	18	290	2	370	
6063	33	0	2065	1	1	26	241	2	-1	
20071	50	0	592	0	0	8	445	4	-1	
20779	50	0	36	0	0	13	104	10	-1	
32548	34	0	703	0	0	17	282	2	150	

	previous	...	month_jun	month_mar	month_may	month_nov	month_oct	\
10547	0	...	1	0	0	0	0	0
40059	2	...	1	0	0	0	0	0
26737	1	...	0	0	0	1	0	0
21071	0	...	0	0	0	0	0	0
39637	0	...	0	0	1	0	0	0
...
38989	1	...	0	0	1	0	0	0
6063	0	...	0	0	1	0	0	0
20071	0	...	0	0	0	0	0	0
20779	0	...	0	0	0	0	0	0
32548	2	...	0	0	0	0	0	0

	month_sep	poutcome_failure	poutcome_other	poutcome_success	\
10547	0	0	0	0	
40059	0	0	0	1	
26737	0	1	0	0	
21071	0	0	0	0	
39637	0	0	0	0	
...	
38989	0	1	0	0	
6063	0	0	0	0	
20071	0	0	0	0	
20779	0	0	0	0	
32548	0	1	0	0	

```

poutcome_unknown
10547      1
40059      0
26737      0
21071      1
39637      1
...
38989      0
6063       1
20071      1
20779      1
32548      0

[13564 rows x 48 columns],
44745      1
44864      0
34749      0
18099      0
1622       0
..
8120       0
39381      0
18512      0
29574      0
655        0
Name: y, Length: 31647, dtype: int32,
10547      0
40059      0
26737      0
21071      0
39637      0
..
38989      0
6063       0
20071      0
20779      0
32548      0
Name: y, Length: 13564, dtype: int32)

```

```
In [22]: from sklearn.linear_model import LogisticRegression
```

```
In [23]: logistic_model = LogisticRegression()
```

```
In [24]: logistic_model.fit(X_train,y_train)
```

```

C:\Users\rajesh\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```

```
Out[24]: LogisticRegression()
```

```
In [25]: # for train data
y_pred = logistic_model.predict(X_train)

# for test data
y_pred_test = logistic_model.predict(X_test)
```

```
In [26]: from sklearn.metrics import confusion_matrix,accuracy_score,classification_report
```

```
In [27]: # for train data
confusion_matrix(y_train,y_pred)
```

```
Out[27]: array([[27607,   381],
               [ 3083,   576]], dtype=int64)
```

```
In [28]: confusion_matrix(y_test,y_pred_test)
```

```
Out[28]: array([[11788, 146],
               [ 1394, 236]], dtype=int64)
```

```
In [29]: #The model accuracy is calculated by (a+d)/(a+b+c+d)
# for train data
print((31340+868)/(31340+590+3370+868))
print(accuracy_score(y_train,y_pred))

0.8905109489051095
0.890542547476854
```

```
In [30]: # for test data
print(accuracy_score(y_test,y_pred_test))

0.8864641698613979
```

```
In [31]: # for train data
print(classification_report(y_train,y_pred))
```

	precision	recall	f1-score	support
0	0.90	0.99	0.94	27988
1	0.60	0.16	0.25	3659
accuracy			0.89	31647
macro avg	0.75	0.57	0.60	31647
weighted avg	0.87	0.89	0.86	31647

```
In [32]: # for test data
print(classification_report(y_test,y_pred_test))
```

	precision	recall	f1-score	support
0	0.89	0.99	0.94	11934
1	0.62	0.14	0.23	1630
accuracy			0.89	13564
macro avg	0.76	0.57	0.59	13564
weighted avg	0.86	0.89	0.85	13564

```
In [33]: from pickle import dump,load
```

```
In [34]: logistic_model.predict(X_test)
```

```
Out[34]: array([0, 0, 0, ..., 0, 0, 0])
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```